Request lifecycle & completion handoff (one operation)

Latency is measured end-to-end, from intended send time to completion — so it includes queueing delay under load

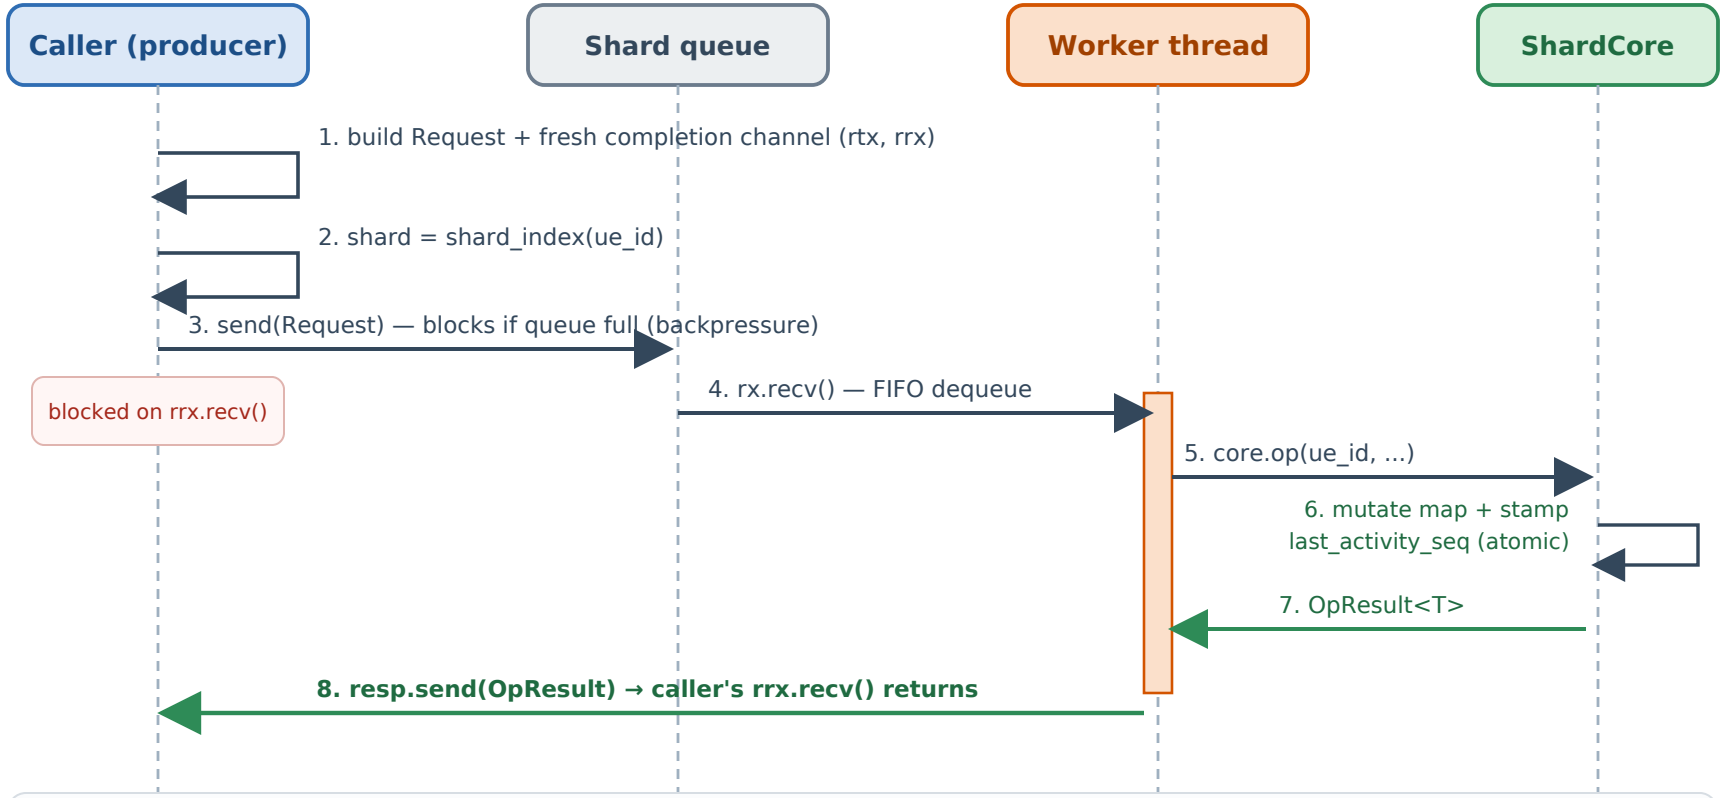

Ownership transfer: the Request (and any heap data it carries) is moved to the worker on submit and is never shared — this upholds the shared-nothing invariant end to end. Requests to the same shard are processed strictly FIFO (O1).

Source: rust/src/service.rs (create_context ... worker_loop) and cpp equivalent (std::promise / std::future).