# Which conclusions the measurement host can carry

A result is final if it is a property of the compilers and the source, and preliminary if it is a property of the machine

## FINAL

Unchanged by the host it was measured on

**Transcript parity**
A conformance suite diffed line by line, not asserted.

**Defect catalogue (RQ2)**
Whether a program compiles is a fact about the compiler.

**Seam structure (RQ4)**
Entry points, unsafe sites, and lines: read off the code.

**Size and complexity proxies**
Counted from the code rather than from a run.

## PRELIMINARY

Carries a contribution from the hypervisor

**Tail latency, p99 and p99.9 (RQ1)**
Computed from exactly the operations a stall disturbs.

**Onset of the overload cliff**
Where saturation begins depends on the scheduler.

**CPU utilisation (RQ3)**
Reverses sign between allocators, reported as unresolved.

**Peak resident set size (RQ3)**
A run-level statistic, so pairing rescues it.

### What the bare-metal protocol buys, and only this

Isolated cores (taskset) | performance governor | Turbo and SMT disabled | Hardware counters

bench/TESTBED.md is scripted and published, and the analysis pipeline does not change: executing the protocol and rerunning the same scripts is the only step between the two columns.

Every row on which the two columns agree is a property of the source code. Every row on which they differ is a property of the machine.