From research question to verdict, grouped by hypothesis

Grouping by hypothesis rather than by question number exposes two things a row-per-question table can only note in passing

QUESTION	OBJECTIVE	HYPOTHESIS	EVIDENCE
RQ1	O1 Quantify the latency distributions to overload, and locate any divergence	**H1** Performance-equivalent below saturation, with any difference confined to the far tail.	§7.5.1 and §7.5.2 — equivalent at the median, unresolvable at the tail
RQ3	O3 Measure the resource cost of the safety model under matched conditions		§7.5.4 — peak RSS reproduces, CPU does not
RQ2	O2 Determine which defect classes are prevented at compile time, and when the rest are caught	**H2** Safe Rust prevents what C++ leaves to testing.	§7.5.5 — five prevented, two compile and misbehave
RQ4	O4 Build and measure an incremental-adoption prototype at the language boundary	**H3** Interop cost is negligible against the work invoked.	§7.5.6 — bounded at 1.37% of the operation

O5 Artifacts and guidance cross-cutting: discharged by the published artifacts and the guidelines, not by any single question.

H1 does double duty for RQ1 and RQ3, so a single wording had to cover equivalence in both its latency and its resource sense — Chapter 8 finds that to be its main weakness.