# From research question to verdict, grouped by hypothesis

Grouping by hypothesis rather than by question number exposes two things a row-per-question table can only note in passing

| QUESTION | OBJECTIVE | HYPOTHESIS | EVIDENCE |
|---|---|---|---|
| **RQ1** | O1 Quantify the latency distributions to overload, and locate any divergence | **H1** Performance-equivalent below saturation, with any difference confined to the far tail. | §7.5.1 and §7.5.2: equivalent at the median, unresolvable at the tail |
| **RQ3** | O3 Measure the resource cost of the safety model under matched conditions | | §7.5.4: peak RSS reproduces, CPU does not |
| **RQ2** | O2 Determine which defect classes are prevented at compile time, and when the rest are caught | **H2** Safe Rust prevents what C++ leaves to testing. | §7.5.5: five prevented, two compile and misbehave |
| **RQ4** | O4 Build and measure an incremental-adoption prototype at the language boundary | **H3** Interop cost is negligible against the work invoked. | §7.5.6: bounded at 10.91% of the operation |

**O5 Artifacts and guidance** cross-cutting: discharged by the published artifacts and the guidelines, not by any single question.

H1 does double duty for RQ1 and RQ3, so a single wording had to cover equivalence in both its latency and its resource sense Chapter 8 finds that to be its main weakness.