Threats to validity, and what each control leaves behind

Shaded by what remains after the control rather than by how the threat was described

	THREAT	CONTROL APPLIED	WHAT REMAINS
INTERNAL Is the language really the only difference?	Implementation bias between the two versions	Frozen specification, shared allocator and optimisation level, pinned hash functions, forced-idiom log, transcript parity	Single-author skill asymmetry. Parity bounds behaviour, not performance
CONSTRUCT Do the metrics measure what the questions ask?	The sweep measures the design, not the language	The operation is measured apart from the handoff and both are reported	One operation on one data structure is narrower than the sweep it supplements
	A defect catalogue is a sample the author chose	Two classes drawn from outside the guarantee, with expectations recorded in advance	Seven programs remain seven programs
	RQ3 captures resource cost but not its origin	None available on this host	The memory result is established but unexplained
	RQ4 measures the seam, not the safety of the whole	The distinction is restated wherever the result is quoted	One address space: a compile-time property of the Rust side bounds nothing at run time
CONCLUSION Are the differences real, or noise?	Run-to-run non-determinism	20 interleaved repeats, bootstrap intervals, an exact paired permutation test, Holm correction, floors quoted	Pairing cancels only what a block shared, so the tail percentiles stay unresolved
	A corrected p-value that does not replicate	The allocator control runs every comparison twice; replication is followed over the p-value	Costs a result where the two disagree: the CPU effect is reported as unresolved
EXTERNAL How far do the results travel?	One subsystem, one workload, one machine	RQ2 and RQ4 rest on the compilers and the C ABI rather than on the host	The performance results do not travel: shared mutable state or allocation churn could behave differently

Two of these are load-bearing

The study cannot rule out single-author skill asymmetry, and its end-to-end instrument answers the question a vendor sizing a node would ask rather than the one a language comparison asks.

Only the last column is a live risk; the middle column is what has already been done about each.
Two rows carry no mitigation at all, and are recorded as such.